

```

In [ ]:
!pip install wandb
import wandb
wandb.init()

Collecting wandb
  Downloading https://files.pythonhosted.org/packages/6c/48/b199e2b3b341ac842108c5db4956091dd75d961cfa77aceb033e99cac20f/wandb-0.10.31-py2.py3-none-any.whl (1.8MB)
    |████████████████████████████████████████| 1.8MB 14.5MB/s
Collecting sentry-sdk<=0.4.0
  Downloading https://files.pythonhosted.org/packages/1c/4a/a54b254f67d8f4052338d54ebe90126f200693440a93ef76d254d581e3ec/sentry_sdk-1.1.0-py2.py3-none-any.whl (131kB)
    |████████████████████████████████████████| 133kB 61.6MB/s
Requirement already satisfied: protobuf>=3.12.0 in /usr/local/lib/python3.7/dist-packages (from wandb) (3.12.4)
Requirement already satisfied: PyYAML in /usr/local/lib/python3.7/dist-packages (from wandb) (3.13)
Requirement already satisfied: psutil>=5.0.0 in /usr/local/lib/python3.7/dist-packages (from wandb) (5.4.8)
Collecting shortuuid>=0.5.0
  Downloading https://files.pythonhosted.org/packages/25/a6/2ecc1daa6a304e7f1b216f0896b26156b78e7c38e1211e9b798b4716c53d/shortuuid-1.0.1-py3-none-any.whl
Requirement already satisfied: Click>=7.0 in /usr/local/lib/python3.7/dist-packages (from wandb) (7.1.2)
Requirement already satisfied: requests<3,>=2.0.0 in /usr/local/lib/python3.7/dist-packages (from wandb) (2.23.0)
Collecting docker-pycreds>=0.4.0
  Downloading https://files.pythonhosted.org/packages/f5/e8/f6bd1eee09314e7e6dee49cb e2c5e22314ccdb38db16c9fc72d2fa80d054/docker_pycreds-0.4.0-py2.py3-none-any.whl
Collecting pathtools
  Downloading https://files.pythonhosted.org/packages/e7/7f/470d6fcdf23f9f3518f6b0b76be9df16dcc8630ad409947f8be2eb0ed13a/pathtools-0.1.2.tar.gz
Collecting GitPython>=1.0.0
  Downloading https://files.pythonhosted.org/packages/27/da/6f6224fdcf47dab57881fe20c0d1bc3122be290198ba0bf26a953a045d92/GitPython-3.1.17-py3-none-any.whl (166kB)
    |████████████████████████████████████████| 174kB 56.4MB/s
Collecting subprocess32>=3.5.3
  Downloading https://files.pythonhosted.org/packages/32/c8/564be4d12629b912ea431f1a50eb8b3b9d00f1a0b1ceff17f266be190007/subprocess32-3.5.4.tar.gz (97kB)
    |████████████████████████████████████████| 102kB 14.5MB/s
Collecting configparser>=3.8.1
  Downloading https://files.pythonhosted.org/packages/fd/01/ff260a18caaf4457eb028c96eeb405c4a230ca06c8ec9c1379f813caa52e/configparser-5.0.2-py3-none-any.whl
Requirement already satisfied: six>=1.13.0 in /usr/local/lib/python3.7/dist-packages (from wandb) (1.15.0)
Requirement already satisfied: python-dateutil>=2.6.1 in /usr/local/lib/python3.7/dist-packages (from wandb) (2.8.1)
Requirement already satisfied: promise<3,>=2.0 in /usr/local/lib/python3.7/dist-packages (from wandb) (2.3)
Requirement already satisfied: urllib3>=1.10.0 in /usr/local/lib/python3.7/dist-packages (from sentry-sdk<=0.4.0->wandb) (1.24.3)
Requirement already satisfied: certifi in /usr/local/lib/python3.7/dist-packages (from sentry-sdk<=0.4.0->wandb) (2020.12.5)
Requirement already satisfied: setuptools in /usr/local/lib/python3.7/dist-packages (from protobuf>=3.12.0->wandb) (56.1.0)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.0.0->wandb) (3.0.4)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.7/dist-packages (from requests<3,>=2.0.0->wandb) (2.10)

```

```

In []: import os          #Import required modules
import cv2
import math
import random
import numpy as np
import datetime as dt
import tensorflow as tf

import matplotlib.pyplot as plt
%matplotlib inline
from sklearn.model_selection import train_test_split
from tensorflow.keras.layers import *
from tensorflow.keras.models import Sequential
from tensorflow.keras.utils import to_categorical
from tensorflow.keras.callbacks import EarlyStopping
from tensorflow.keras.utils import plot_model

In []: seed_constant = 23  #Initialize random number generator
np.random.seed(seed_constant)
random.seed(seed_constant)
tf.random.set_seed(seed_constant)

```

```

In []: from google.colab import drive
drive.mount('/gdrive')
%cd /gdrive

```

Mounted at /gdrive
/gdrive

```

In []: %cd 'My Drive/'

```

```

%cd 'Action Recognition'

```

/gdrive/My Drive
/gdrive/My Drive/Action Recognition

```

In []: %ls

```

```

hmdb51/
hmdb512/
hmdb51_org.rar
hmdb51rars/
KTH/

```

```

Model__Date_Time_2021_05_11__18_40_06__Loss_0.46230462193489075__Accuracy_0.99260
41960716248.h5
Model__Date_Time_2021_05_11__19_57_16__Loss_0.003302493831142783__Accuracy_0.9992
708563804626.h5
Model__Date_Time_2021_05_11__23_01_06__Loss_0.001863537821918726__Accuracy_0.9994
791746139526.h5
Model__Date_Time_2021_05_11__23_34_47__Loss_0.005838708486407995__Accuracy_0.9990
624785423279.h5
Model__Date_Time_2021_05_14__18_33_00__Loss_0.005261874292045832__Accuracy_0.9986
458420753479.h5
Model__Date_Time_2021_05_14__19_02_48__Loss_0.0034059989266097546__Accuracy_0.999
1666674613953.h5

```

```

In [ ]:
    #Constants

    datasetName= 'hmdb51'                #Chosen Dataset

In...
    image_height, image_width = 64, 64      #Set pixel values
    images_per_class = 8000                 #Set number of frames from each video class
    dataset_directory = "hmdb51"           #Set dataset name
    classes_list = ["pullup", "punch", "dive", "fencing", "ride_bike", "golf"] #Cho
    model_output_size = len(classes_list)   #test

In...
    def frames_extraction(video_path): #helper function to extract frames from video
        frames_list = []               #empty list for frames
        video_reader = cv2.VideoCapture(video_path) #Read frames from video
        while True:                    #Iterate through frames
            success, frame = video_reader.read() #Whilst frames are available
            if not success:
                break
            resized_frame = cv2.resize(frame, (image_height, image_width)) #Resize
            normalized_frame = resized_frame / 255                          #Normal
            frames_list.append(normalized_frame)                            #Add to
        video_reader.release()                                              #Close
        return frames_list                                                 #Retur

l...
    def create_dataset(): #Create dataset function
        temp_features = [] #empty list to hold each videos frames
        features = [] #final list of frames will be in this list
        labels = [] #Final list of labels will be in this list

        for class_index, class_name in enumerate(classes_list): #Iterate through cho:
            print(f'Extracting Data of Class: {class_name}')
            files_list = os.listdir(os.path.join(dataset_directory, class_name)) #Got
            for file_name in files_list:
                video_file_path = os.path.join(dataset_directory, class_name, file_name)
                frames = frames_extraction(video_file_path) #Extract frames for curri
                temp_features.extend(frames) #Add to temp frames list

            features.extend(random.sample(temp_features, images_per_class)) #Cho
            labels.extend([class_index] * images_per_class) #
            temp_features.clear()

        features = np.asarray(features) #Convert both to numpy array
        labels = np.array(labels)
        return features, labels

In [ ]:
    features, labels = create_dataset() #Fetch data

Extracting Data of Class: pullup
Extracting Data of Class: punch
Extracting Data of Class: dive
Extracting Data of Class: fencing
Extracting Data of Class: ride_bike
Extracting Data of Class: golf
In [ ]:

```

```

(48000, 64, 64, 3)
(48000,)
In ...
    one_hot_encoded_labels = to_categorical(labels) #convert labels into one-hot-enc

l...
    features_train, features_test, labels_train, labels_test = train_test_split(featur

l...
def create_model(): #Create NN
    model = Sequential() #Keras sequential model

    model.add(Conv2D(filters = 64, kernel_size = (3, 3), activation = 'relu', inpu
    model.add(Conv2D(filters = 64, kernel_size = (3, 3), activation = 'relu'))
    model.add(Conv2D(filters = 64, kernel_size = (3, 3), activation = 'relu'))
    model.add(Conv2D(filters = 64, kernel_size = (3, 3), activation = 'relu'))
    model.add(Conv2D(filters = 64, kernel_size = (3, 3), activation = 'relu'))
    model.add(BatchNormalization())
    model.add(MaxPooling2D(pool_size = (2, 2)))
    model.add(GlobalAveragePooling2D())
    model.add(Dense(288, activation = 'relu'))
    model.add(Dense(288, activation = 'relu'))
    model.add(BatchNormalization())
    model.add(Dense(model_output_size, activation = 'softmax'))

    model.summary() #Show model summary

    return model

model = create_model()

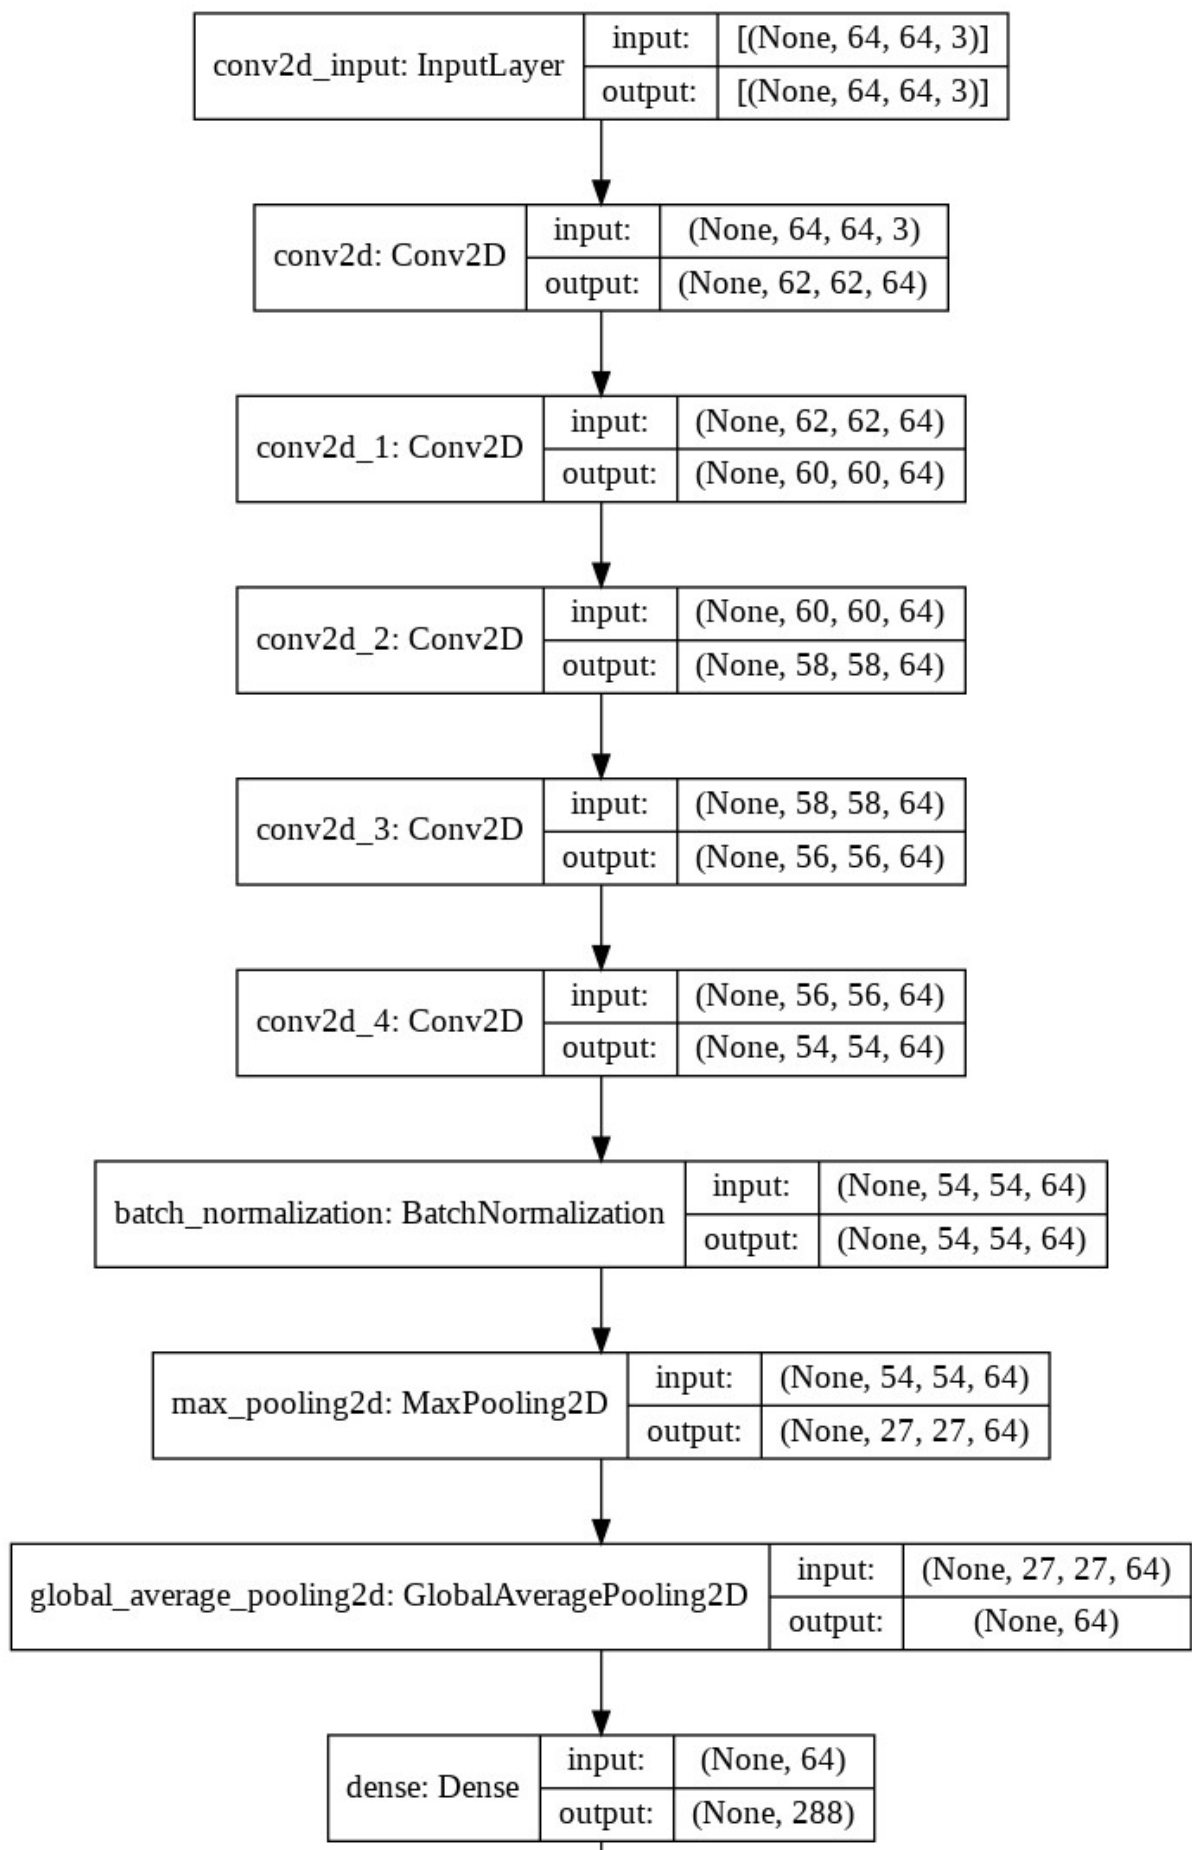
```

Model: "sequential"

Layer (type)	Output Shape	Param #
=====		
conv2d (Conv2D)	(None, 62, 62, 64)	1792
conv2d_1 (Conv2D)	(None, 60, 60, 64)	36928
conv2d_2 (Conv2D)	(None, 58, 58, 64)	36928
conv2d_3 (Conv2D)	(None, 56, 56, 64)	36928
conv2d_4 (Conv2D)	(None, 54, 54, 64)	36928
batch_normalization (BatchNo	(None, 54, 54, 64)	256
max_pooling2d (MaxPooling2D)	(None, 27, 27, 64)	0
global_average_pooling2d (Gl	(None, 64)	0
dense (Dense)	(None, 288)	18720
dense_1 (Dense)	(None, 288)	83232
batch_normalization_1 (Batch	(None, 288)	1152

In ... `plot_model(model, show_shapes = True, show_layer_names = True) #Plot model diagr`

Out...



```

l...
from keras import optimizers
import keras
optimizer = keras.optimizers.Adam(lr=0.0001)

early_stopping_callback = EarlyStopping(monitor = 'val_loss', patience = 10, mode

model.compile(loss = 'categorical_crossentropy', optimizer = optimizer, metrics =

model_training_history = model.fit(x = features_train, y = labels_train, epochs =

/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/optimizer_v2/optimize
r_v2.py:375: UserWarning: The `lr` argument is deprecated, use `learning_rate` inste
ad.
"The `lr` argument is deprecated, use `learning_rate` instead.")
Epoch 1/40
1920/1920 [=====] - 34s 10ms/step - loss: 0.5631 - accurac
y: 0.8086 - val_loss: 0.2843 - val_accuracy: 0.9128
Epoch 2/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.2154 - accurac
y: 0.9322 - val_loss: 0.2088 - val_accuracy: 0.9211
Epoch 3/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.1217 - accurac
y: 0.9627 - val_loss: 0.0928 - val_accuracy: 0.9737
Epoch 4/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0906 - accurac
y: 0.9722 - val_loss: 0.1249 - val_accuracy: 0.9583
Epoch 5/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0682 - accurac
y: 0.9801 - val_loss: 0.1521 - val_accuracy: 0.9424
Epoch 6/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0531 - accurac
y: 0.9838 - val_loss: 0.0269 - val_accuracy: 0.9924
Epoch 7/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0420 - accurac
y: 0.9877 - val_loss: 0.0351 - val_accuracy: 0.9904
Epoch 8/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0408 - accurac
y: 0.9876 - val_loss: 0.0758 - val_accuracy: 0.9771
Epoch 9/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0311 - accurac
y: 0.9906 - val_loss: 0.1535 - val_accuracy: 0.9516
Epoch 10/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0333 - accurac
y: 0.9898 - val_loss: 0.1692 - val_accuracy: 0.9430
Epoch 11/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0289 - accurac
y: 0.9909 - val_loss: 0.0062 - val_accuracy: 0.9984
Epoch 12/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0232 - accurac
y: 0.9931 - val_loss: 0.0125 - val_accuracy: 0.9956
Epoch 13/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0272 - accurac
y: 0.9916 - val_loss: 0.0845 - val_accuracy: 0.9728
Epoch 14/40
1920/1920 [=====] - 19s 10ms/step - loss: 0.0174 - accurac
y: 0.9950 - val_loss: 0.0140 - val_accuracy: 0.9964
Epoch 15/40

```

```

In[: model_evaluation_history = model.evaluate(features_test, labels_test)

from sklearn.metrics import classification_report

y_pred = model.predict(features_test, batch_size=4, verbose=1)
y_pred_bool = np.argmax(y_pred, axis=1)

l_test=np.argmax(labels_test, axis=1)

print(classification_report(l_test, y_pred_bool))

300/300 [=====] - 2s 5ms/step - loss: 0.0035 - accuracy: 0.9986
2400/2400 [=====] - 4s 1ms/step
      precision    recall  f1-score   support

     0       1.00      1.00      1.00     1604
     1       1.00      1.00      1.00     1615
     2       1.00      1.00      1.00     1593
     3       1.00      1.00      1.00     1585
     4       1.00      1.00      1.00     1602
     5       1.00      1.00      1.00     1601

 accuracy                   1.00      9600
 macro avg       1.00      1.00      1.00      9600
 weighted avg    1.00      1.00      1.00      9600

In[: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(l_test, y_pred_bool)

print (cm)

[[1604    0    0    0    0    0]
 [   0 1609    1    2    3    0]
 [   0    1 1589    1    2    0]
 [   0    1    2 1582    0    0]
 [   0    0    0    0 1602    0]
 [   0    0    0    0    0 1601]]

l...
# Creating a useful name for our model, incase you're saving multiple models (OPT.
date_time_format = '%Y_%m_%d_%H_%M_%S'
current_date_time_dt = dt.datetime.now()
current_date_time_string = dt.datetime.strftime(current_date_time_dt, date_time_f
model_evaluation_loss, model_evaluation_accuracy = model_evaluation_history
model_name = f'Model__Date_Time_{current_date_time_string}__Loss_{model_evaluat


# Saving your Model
model.save(model_name)

In[: def plot_metric(metric_name_1, metric_name_2, plot_name):
    # Get Metric values using metric names as identifiers
    metric_value_1 = model_training_history.history[metric_name_1]
    metric_value_2 = model_training_history.history[metric_name_2]


    # Constructing a range object which will be used as time

```

```
In [ ]: plot_metric('loss', 'val_loss', 'Total Loss vs Total Validation Loss')
```



```
In ... plot_metric('accuracy', 'val_accuracy', 'Total Accuracy vs Total Validation Accu')
```



```
In [ ]: from collections import Counter

def get_first_mode(a):
    c = Counter(a)
    mode_count = max(c.values())
    mode = {key for key, count in c.items() if count == mode_count}
    first_mode = next(x for x in a if x in mode)
    return first_mode
```

```
In [ ]: def frames_extraction2(video_path):
    frames_list = []

    vidObj = cv2.VideoCapture(video_path)
```

Used as counter variable

In ...

```
#Evaluating a different dataset
```

```
from tqdm import tqdm
from statistics import mode
```

```
predict = []
actual = []
dataset_directory2="UCF50"
```

```
# Declaring Empty Lists to store the features and labels values.
temp_features = []
features = []
labels = []
```

```
cc=0
```

```
# Iterating through all the classes mentioned in the classes list
for class_index, class_name in enumerate(classes_list):
    print(f'Extracting Data of Class: {class_name}')
```

```
# Getting the list of video files present in the specific class name directory
files_list = os.listdir(os.path.join(dataset_directory2, class_name))
```

```
# Iterating through all the files present in the files list
for file_name in files_list:
```

```
    # Construct the complete video path
    video_file_path = os.path.join(dataset_directory2, class_name, file_name)
```

```
    # Calling the frame_extraction method for every video file path
    frames = frames_extraction2(video_file_path)
```

```
    temppred=[]
```

```
    for i in frames:
        temppred.append(model.predict_classes(np.expand_dims(i, axis = 0))[0])
```

```
    print (temppred)
    print ("mode", get_first_mode(temppred), cc)
    cc+=1
    predict.append(get_first_mode(temppred))
    actual.append(class_index)
```

Extracting Data of Class: pullup

Defected frame

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/sequential.py:
455: UserWarning: `model.predict_classes()` is deprecated and will be removed after
2021-01-01. Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your m
odel does multi-class classification (e.g. if it uses a `softmax` last-layer activ
ation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary cla
ssification (e.g. if it uses a `sigmoid` last-layer activation).
```

```
warnings.warn('`model.predict_classes()` is deprecated and '
```

```
[1, 1, 1, 1, 1, 1, 1]
```

mode 1 0

Defected frame

```
In []: print(classification_report(actual, predict))
```

	precision	recall	f1-score	support
0	0.89	0.39	0.54	120
1	0.44	0.34	0.39	160
2	0.51	0.75	0.60	153
3	0.30	0.22	0.25	111
4	0.43	0.71	0.54	145
5	0.87	0.68	0.77	142
accuracy			0.53	831
macro avg	0.57	0.52	0.52	831
weighted avg	0.57	0.53	0.52	831

```
In []: print(confusion_matrix(actual, predict))
```

```
[[ 47  32   9   8  15   9]
 [  4  55  38  13  50   0]
 [  2   6 114  13  16   2]
 [  0  21  25  24  41   0]
 [  0   5  16  18 103   3]
 [  0   5  22   4  14  97]]
```