

```
In []:
    #!/pip install wandb
    #import wandb
    #wandb.init()
```

```
In []:
    import os
    import cv2
    import math
    import random
    import numpy as np
    import datetime as dt
    import tensorflow as tf
    from tensorflow import keras

    import matplotlib.pyplot as plt
    %matplotlib inline

    from sklearn.model_selection import train_test_split

    from tensorflow.keras.layers import *
    from tensorflow.keras.models import Sequential
    from tensorflow.keras.utils import to_categorical
    from tensorflow.keras.callbacks import EarlyStopping
    from tensorflow.keras.utils import plot_model

    from keras.applications.vgg16 import VGG16
    from keras.applications.vgg16 import preprocess_input
```

```
In []:
    from google.colab import drive
    drive.mount('/gdrive')
    %cd /gdrive
```

Drive already mounted at /gdrive; to attempt to forcibly remount, call drive.mount("/gdrive", force_remount=True).

/gdrive

```
In []:
    %cd 'My Drive'
```

```
    %cd 'Action Recognition'
```

```
/gdrive/My Drive
/gdrive/My Drive/Action Recognition
```

```
In []:
    image_height, image_width = 64, 64
    images_per_class = 8000
    dataset_directory = "hmdb51"
    classes_list = ["pullup", "punch", "dive", "fencing", "ride_bike", "golf"]
    model_output_size = len(classes_list)
```

```
In []:
    def frames_extraction(video_path):
        frames_list = []
        video_reader = cv2.VideoCapture(video_path)
```

```
def create_dataset():
```

```
features = np.asarray(features)
labels = np.array(labels)
```

```
features, labels = create_dataset()
```

[illegible]

```

In [: seed_constant = 23
      np.random.seed(seed_constant)
      random.seed(seed_constant)
      tf.random.set_seed(seed_constant)

l...
print (features.shape)
print (labels.shape)

one_hot_encoded_labels = to_categorical(labels)

features_train, features_test, labels_train, labels_test = train_test_split(features, labels,
                                     test_size=0.2, random_state=seed_constant)

print (features_train.shape)
print (labels_train.shape)

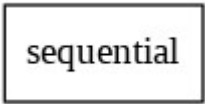
(48000, 64, 64, 3)
(48000,)
(38400, 64, 64, 3)
(38400, 6)
In [: print (labels_train)

[[0. 0. 1. 0. 0. 0.]
 [1. 0. 0. 0. 0. 0.]
 [0. 0. 0. 0. 1. 0.]
 ...
 [0. 0. 0. 1. 0. 0.]
 [0. 1. 0. 0. 0. 0.]
 [0. 0. 0. 0. 0. 1.]]
In...
# Load model
base_model = Sequential()
base_model.add(VGG16(input_shape=(64,64,3), weights='imagenet', include_top=False))
base_model.add(Dense(288, activation = 'relu'))
base_model.add(Dense(288, activation = 'relu'))
base_model.add(Dense(6, activation='softmax'))
# summarize the model

base_model.layers[0].trainable = False

In ...
base_model.compile(optimizer='sgd',loss='categorical_crossentropy',metrics=['acc

In [: plot_model(base_model,show_shapes = True, show_layer_names = True)

Out[: 

l...
# Adding Early Stopping Callback
early_stopping_callback = EarlyStopping(monitor = 'val_loss', patience = 10, mode

# Start Training
model_training_history = base_model.fit(x = features_train, y = labels_train, epoc

```

Epoch 1/40
1920/1920 [=====] - 35s 10ms/step - loss: 0.1594 - accuracy: 0.9585 - val_loss: 0.0309 - val_accuracy: 0.9908

Epoch 2/40
1920/1920 [=====] - 18s 9ms/step - loss: 0.0101 - accuracy: 0.9976 - val_loss: 0.0147 - val_accuracy: 0.9960

Epoch 3/40
1920/1920 [=====] - 18s 10ms/step - loss: 0.0030 - accuracy: 0.9998 - val_loss: 0.0113 - val_accuracy: 0.9966

Epoch 4/40
1920/1920 [=====] - 18s 10ms/step - loss: 0.0015 - accuracy: 1.0000 - val_loss: 0.0096 - val_accuracy: 0.9975

Epoch 5/40
1920/1920 [=====] - 18s 10ms/step - loss: 9.9968e-04 - accuracy: 1.0000 - val_loss: 0.0088 - val_accuracy: 0.9975

Epoch 6/40
1920/1920 [=====] - 18s 9ms/step - loss: 7.3064e-04 - accuracy: 1.0000 - val_loss: 0.0080 - val_accuracy: 0.9979

Epoch 7/40
1920/1920 [=====] - 18s 9ms/step - loss: 5.8856e-04 - accuracy: 1.0000 - val_loss: 0.0077 - val_accuracy: 0.9978

Epoch 8/40
1920/1920 [=====] - 18s 10ms/step - loss: 4.8512e-04 - accuracy: 1.0000 - val_loss: 0.0077 - val_accuracy: 0.9978

Epoch 9/40
1920/1920 [=====] - 18s 9ms/step - loss: 4.1695e-04 - accuracy: 1.0000 - val_loss: 0.0075 - val_accuracy: 0.9977

Epoch 10/40
1920/1920 [=====] - 18s 9ms/step - loss: 3.6671e-04 - accuracy: 1.0000 - val_loss: 0.0071 - val_accuracy: 0.9978

Epoch 11/40
1920/1920 [=====] - 18s 9ms/step - loss: 3.2759e-04 - accuracy: 1.0000 - val_loss: 0.0070 - val_accuracy: 0.9979

Epoch 12/40
1920/1920 [=====] - 18s 9ms/step - loss: 2.9161e-04 - accuracy: 1.0000 - val_loss: 0.0071 - val_accuracy: 0.9978

Epoch 13/40
1920/1920 [=====] - 18s 9ms/step - loss: 2.6644e-04 - accuracy: 1.0000 - val_loss: 0.0070 - val_accuracy: 0.9980

Epoch 14/40
1920/1920 [=====] - 18s 10ms/step - loss: 2.4417e-04 - accuracy: 1.0000 - val_loss: 0.0068 - val_accuracy: 0.9979

Epoch 15/40
1920/1920 [=====] - 18s 9ms/step - loss: 2.2375e-04 - accuracy: 1.0000 - val_loss: 0.0069 - val_accuracy: 0.9980

Epoch 16/40
1920/1920 [=====] - 18s 9ms/step - loss: 2.0753e-04 - accuracy: 1.0000 - val_loss: 0.0067 - val_accuracy: 0.9979

Epoch 17/40
1920/1920 [=====] - 18s 9ms/step - loss: 1.9307e-04 - accuracy: 1.0000 - val_loss: 0.0068 - val_accuracy: 0.9980

Epoch 18/40
1920/1920 [=====] - 18s 10ms/step - loss: 1.8017e-04 - accuracy: 1.0000 - val_loss: 0.0067 - val_accuracy: 0.9979

Epoch 19/40
1920/1920 [=====] - 18s 9ms/step - loss: 1.6935e-04 - accuracy: 1.0000 - val_loss: 0.0066 - val_accuracy: 0.9980

Epoch 20/40
1920/1920 [=====] - 18s 9ms/step - loss: 1.5937e-04 - accuracy:

```
In[: model_evaluation_history = base_model.evaluate(features_test, labels_test)

from sklearn.metrics import classification_report

y_pred = base_model.predict(features_test, batch_size=4, verbose=1)
y_pred_bool = np.argmax(y_pred, axis=1)

l_test=np.argmax(labels_test, axis=1)

print(classification_report(l_test, y_pred_bool))
```

300/300 [=====] - 4s 11ms/step - loss: 0.0114 - accuracy: 0.9972

2400/2400 [=====] - 9s 4ms/step

	precision	recall	f1-score	support
0	1.00	1.00	1.00	1604
1	1.00	1.00	1.00	1615
2	1.00	0.99	1.00	1593
3	0.99	1.00	0.99	1585
4	1.00	1.00	1.00	1602
5	1.00	1.00	1.00	1601
accuracy			1.00	9600
macro avg	1.00	1.00	1.00	9600
weighted avg	1.00	1.00	1.00	9600

```
In[: from sklearn.metrics import confusion_matrix
cm = confusion_matrix(l_test, y_pred_bool)

print (cm)
```

```
[[1604    0    0    0    0    0]
 [   1 1607    2    5    0    0]
 [   0    3 1585    2    1    2]
 [   0    5    2 1578    0    0]
 [   0    0    2    2 1598    0]
 [   0    0    0    0    0 1601]]
```

```
In[: def plot_metric(metric_name_1, metric_name_2, plot_name):
    # Get Metric values using metric names as identifiers
    metric_value_1 = model_training_history.history[metric_name_1]
    metric_value_2 = model_training_history.history[metric_name_2]


    # Constructing a range object which will be used as time
    epochs = range(len(metric_value_1))

    # Plotting the Graph
    plt.plot(epochs, metric_value_1, 'blue', label = metric_name_1)
    plt.plot(epochs, metric_value_2, 'red', label = metric_name_2)


    # Adding title to the plot
    plt.title(str(plot_name))

    # Adding Legend to the plot
    plt.legend()
```

```
In [ ]: plot_metric('loss', 'val_loss', 'Total Loss vs Total Validation Loss')
```



```
In ... plot_metric('accuracy', 'val_accuracy', 'Total Accuracy vs Total Validation Accu')
```



```
In [ ]: from collections import Counter

def get_first_mode(a):
    c = Counter(a)
    mode_count = max(c.values())
    mode = {key for key, count in c.items() if count == mode_count}
    first_mode = next(x for x in a if x in mode)
    return first_mode
```

```
In [ ]: def frames_extraction2(video_path):
    frames_list = []

    vidObj = cv2.VideoCapture(video_path)
    skip_frames=30
```

Used as counter variable

In ...

```
#Evaluating a different dataset
```

```
from tqdm import tqdm
from statistics import mode
```

```
predict = []
actual = []
dataset_directory2="UCF50"
```

```
temp_features = []
features = []
labels = []
```

```
cc=0
```

```
for class_index, class_name in enumerate(classes_list):
    print(f'Extracting Data of Class: {class_name}')
```

```
    files_list = os.listdir(os.path.join(dataset_directory2, class_name))
```

```
    for file_name in files_list:
```

```
        video_file_path = os.path.join(dataset_directory2, class_name, file_name)
```

```
        frames = frames_extraction2(video_file_path)
```

```
        temppred=[]
```

```
        for i in frames:
```

```
            temppred.append(base_model.predict_classes(np.expand_dims(i, axis = 0))
```

```
            print (temppred)
```

```
            print ("mode", get_first_mode(temppred), cc)
```

```
            cc+=1
```

```
            predict.append(get_first_mode(temppred))
```

```
            actual.append(class_index)
```

Extracting Data of Class: pullup

Defected frame

```
/usr/local/lib/python3.7/dist-packages/tensorflow/python/keras/engine/sequential.py:
455: UserWarning: `model.predict_classes()` is deprecated and will be removed after
2021-01-01. Please use instead: * `np.argmax(model.predict(x), axis=-1)`, if your m
odel does multi-class classification (e.g. if it uses a `softmax` last-layer activ
ation). * `(model.predict(x) > 0.5).astype("int32")`, if your model does binary cla
ssification (e.g. if it uses a `sigmoid` last-layer activation).
```

```
warnings.warn("`model.predict_classes()` is deprecated and '
```

```
[0, 0, 0, 0, 0, 5, 0]
```

```
mode 0 0
```

Defected frame

```
[0, 3, 2, 3, 2, 2]
```

```
mode 2 1
```

Defected frame

```
[2, 2, 2, 2, 2, 2, 2, 1]
```

```
mode 2 2
```

Defected frame

```
In [26]: print(classification_report(actual, predict))
```

	precision	recall	f1-score	support
0	0.91	0.57	0.70	120
1	0.62	0.28	0.39	160
2	0.47	0.96	0.63	153
3	0.31	0.21	0.25	111
4	0.70	0.83	0.76	145
5	0.76	0.66	0.71	142
accuracy			0.60	831
macro avg	0.63	0.59	0.57	831
weighted avg	0.63	0.60	0.58	831

```
In [27]: print(confusion_matrix(actual, predict))
```

```
[[ 68  9 16  9 11  7]
 [  5 45 60 34 14  2]
 [  0  2 147  0  4  0]
 [  1 12 44 23 15 16]
 [  0  1 18  1 121  4]
 [  1  3 29  8  7 94]]
```